Two Graph Coloring Algorithms Based on Independent Sets: IndSet-Based Exact Coloring and IndSet-Removal Coloring

Graph Exact Coloring Heuristic Project

March 30, 2026

Abstract

We describe and analyze two graph coloring algorithms both rooted in the independent-set structure of a graph. The first, *IndSet-Based exact coloring*, determines the chromatic number exactly by evaluating a combinatorial quantity derived from inclusion-exclusion over all independent sets. The quantity, denoted $c_k(G)$, satisfies $c_k(G) = 0$ if and only if $k < \chi(G)$, so the chromatic number equals the smallest k for which $c_k(G) \neq 0$. The second, *IndSet-Removal Coloring*, repeatedly peels off one independent set per round and assigns it the next available color, yielding a fast upper bound on $\chi(G)$. Both algorithms are implemented in C++17 using the Boost Graph Library. We present experimental results on cycles C_n, complete graphs K_n, the Petersen graph $GP(5, 2)$ and other generalized Petersen graphs, wheel graphs W_n, complete multipartite graphs $K_{s,\ldots,s}$, edgeless graphs, and graphs from the SuiteSparse Matrix Collection, reporting both the exact chromatic number and the IndSet-Removal upper bound side by side. The exact algorithm is practical for graphs up to roughly 13–14 vertices on a single core; IndSet-Removal runs in $O(n \cdot m)$ and is applicable to larger instances.

Contents

1 Introduction **2**

2 Preliminaries **3**

3 The Exact Algorithm **3**
- 3.1 The Combinatorial Formula . . . 3
- 3.2 Combinatorial Interpretation . . . 3
- 3.3 Equivalence with k-Colorability . . . 4
- 3.4 Algorithm . . . 4

4 IndSet-Removal Coloring **5**
- 4.1 Independent Set Extraction . . . 5
- 4.2 IndSet-Removal Coloring . . . 5
- 4.3 Quality and Complexity . . . 5

5 Implementation **6**

6 Experimental Results **7**
6.1 Exact Values of c_k at the Chromatic Number . . . 7
6.2 Parity Property of Odd Cycles . . . 7
6.3 Petersen Graph . . . 7
6.4 Wheel Graphs . . . 7
6.5 Complete Multipartite Graphs . . . 8
6.6 Generalized Petersen Family . . . 8
6.7 SuiteSparse Matrix Collection Graphs . . . 8
6.8 Scaling Analysis . . . 8
6.9 High Chromatic Number Graphs . . . 9
6.10 IndSet-Removal vs. Exact Coloring . . . 9

7 Complexity Analysis **9**
7.1 Independent Set Enumeration . . . 9
7.2 Evaluation of c_k . . . 10
7.3 Overall Complexity . . . 10
7.4 Space Complexity . . . 10

8 Conclusion **10**

1 Introduction

Graph coloring is one of the most extensively studied problems in combinatorics and theoretical computer science. Given a graph $G = (V, E)$, a *proper k-coloring* assigns one of k colors to each vertex so that no two adjacent vertices share a color. The minimum such k is the *chromatic number* $\chi(G)$. Determining $\chi(G)$ is NP-hard in general [1], yet exact algorithms are essential for small or sparse graphs arising in scientific computing, scheduling, and register allocation.

This paper documents two algorithms implemented in the `graph-exact-coloring-heuristic` project. The first, *IndSet-Based Exact Coloring*, uses the core function `c_k` which evaluates the inclusion-exclusion sum

$$c_k(G, k) \;=\; \sum_{S \subseteq V} (-1)^{|S|} \left|\{\, I \in \mathcal{F}(G) : I \cap S = \emptyset \,\}\right|^k,$$

where $\mathcal{F}(G)$ is the family of all independent sets of G (including the empty set). We show that this quantity is zero precisely when G is not k-colorable, and hence one can binary-search (or linearly scan) over k to find $\chi(G)$.

The second, *IndSet-Removal Coloring* (implemented in `demo_is_removal_coloring.cpp`), repeatedly extracts an independent set and assigns it the next available color, yielding an upper bound that may exceed $\chi(G)$ in general.

Organization. Section 2 fixes notation and recalls standard facts about graph coloring. Section 3 derives the formula and proves its correctness. Section 4 describes the IndSet-Removal Coloring algorithm. Section 5 outlines the C++ implementation. Section 6 presents experimental results, including graphs from the SuiteSparse Matrix Collection. Section 7 analyzes time and space complexity. Section 8 concludes.

2 Preliminaries

Throughout, $G = (V, E)$ is a simple undirected graph with $n = |V|$ vertices and $m = |E|$ edges. Vertices are labeled $0, 1, \ldots, n-1$.

Definition 1 (Independent set)**.** *A set $I \subseteq V$ is* independent *if no two vertices in I are adjacent. We write $\mathcal{F}(G)$ for the family of all independent sets of G, including $\emptyset$.*

Definition 2 (Proper coloring and chromatic number)**.** *A* proper k-coloring *of G is a function $f : V \to \{1, \ldots, k\}$ with $f(u) \neq f(v)$ whenever $uv \in E$. The* chromatic number $\chi(G)$ *is the minimum k for which a proper k-coloring exists.*

Definition 3 (Independent set cover)**.** *An* ordered k-cover by independent sets *(hereafter, k-IndSet-cover) is a k-tuple $(I_1, \ldots, I_k) \in \mathcal{F}(G)^k$ such that $I_1 \cup \cdots \cup I_k = V$.*

Note that in a k-IndSet-cover the sets need not be disjoint, and they may be empty. A proper k-coloring induces a k-IndSet-cover (the color classes), but a k-IndSet-cover may have overlapping sets. We will show these notions are equivalent for the *existence* question.

Definition 4 (Chromatic polynomial)**.** *The* chromatic polynomial $P(G, k)$ *counts the number of proper k-colorings of G. It satisfies $P(G, k) = 0$ iff $k < \chi(G)$.*

Well-known chromatic polynomials for families used in experiments:

$$\begin{aligned} P(K_n, k) &= k(k-1)(k-2)\cdots(k-n+1), \\ P(C_n, k) &= (k-1)^n + (-1)^n(k-1), \\ P(E_n, k) &= k^n, \end{aligned}$$

where K_n is the complete graph, C_n the cycle, and E_n the edgeless graph on n vertices.

3 The Exact Algorithm

3.1 The Combinatorial Formula

For a graph G and a positive integer k, define

$$\alpha(S, \mathcal{F}) \;=\; \big|\{\, I \in \mathcal{F}(G) : I \cap S = \emptyset \,\}\big|$$

to be the number of independent sets of G that avoid S entirely. Equivalently, $\alpha(S, \mathcal{F})$ is the total number of independent sets in the subgraph of G induced by $V \setminus S$. The algorithm evaluates

$$c_k(G, k) \;=\; \sum_{S \subseteq V} (-1)^{|S|}\, \alpha(S, \mathcal{F})^k. \tag{1}$$

3.2 Combinatorial Interpretation

Theorem 1. *$c_k(G, k)$ equals the number of ordered k-IndSet-covers of G.*

Proof. Let $N = |\mathcal{F}(G)|$. The number of ordered k-tuples $(I_1, \ldots, I_k) \in \mathcal{F}(G)^k$ with no constraint is N^k. For a fixed vertex $v \in V$ let A_v be the set of k-tuples in which v is covered by no I_j (i.e. $v \notin I_j$ for all j). Then

$$|A_v| = |\{I \in \mathcal{F}(G) : v \notin I\}|^k = \alpha(\{v\}, \mathcal{F})^k.$$

More generally, for any $S \subseteq V$,

$$\Big|\bigcap_{v\in S} A_v\Big| = |\{I \in \mathcal{F}(G) : S \cap I = \emptyset\}|^k = \alpha(S,\mathcal{F})^k.$$

By inclusion-exclusion, the number of k-tuples in which every vertex is covered by at least one I_j is

$$\sum_{S\subseteq V} (-1)^{|S|}\, \alpha(S,\mathcal{F})^k \;=\; c_k(G,k). \qquad \square$$

3.3 Equivalence with k-Colorability

Theorem 2. *$c_k(G,k) = 0$ if and only if G is not properly k-colorable, i.e. $k < \chi(G)$.*

Proof. ($\Leftarrow$) Suppose $k \geq \chi(G)$. Fix a proper k-coloring $f : V \to \{1,\ldots,k\}$ and let $I_j = f^{-1}(j)$ for each j. Each I_j is an independent set and $\bigcup_j I_j = V$, so $(I_1,\ldots,I_k)$ is a k-IndSet-cover. Hence $c_k(G,k) \geq 1 > 0$.

($\Rightarrow$) Suppose $k < \chi(G)$. We show no k-IndSet-cover exists. If $(I_1,\ldots,I_k)$ were a k-IndSet-cover, define $f(v) = \min\{j : v \in I_j\}$ for each $v \in V$. Since I_j is independent, no two adjacent vertices can share the same j; thus f is a proper k-coloring, contradicting $k < \chi(G)$. Therefore $c_k(G,k) = 0$. $\square$

Corollary 3. $\chi(G) = \min\{\, k \geq 1 : c_k(G,k) \neq 0 \,\}$.

Remark 1. *The quantity $c_k(G,k)$ is not the chromatic polynomial $P(G,k)$ in general. For $k = \chi(G)$, the two coincide for complete graphs (where every k-IndSet-cover induces a proper coloring immediately), but they differ for most other graphs at values $k > \chi(G)$. For example, $c_k(C_5,3) = 60$ while $P(C_5,3) = 30$. Nevertheless, both are zero for $k < \chi(G)$ and nonzero for $k \geq \chi(G)$, which is the only property needed to determine the chromatic number.*

3.4 Algorithm

The full procedure to compute $\chi(G)$ is given in Algorithm 2. It first enumerates all independent sets of G (Algorithm 1) using a power-set sweep with bitset membership testing, then evaluates formula (1) for increasing k.

Algorithm 1 Enumerate all independent sets of G

Require: Graph $G = (V,E)$ with $|V| = n$
Ensure: Set $\mathcal{F}$ of all independent sets (as n-bit integers)

```
F ← ∅
for i ← 0 to 2^n − 1 do
    B ← n-bit representation of i
    if ∀ u ≠ v: B[u] = B[v] = 1 ⇒ uv ∉ E then
        F ← F ∪ {B}
    end if
end forreturn F
```

Bitset Representation

Each independent set is stored as an n-bit integer (a `boost::dynamic_bitset`). The independence check for a candidate set B iterates over all pairs (i,j) with $B[i] = B[j] = 1$ and verifies $ij \notin E$. The subset-avoidance test $I \cap S = \emptyset$ reduces to the bitwise operation $(I \,\&\, S) = 0$.

Algorithm 2 Compute $\chi(G)$ via inclusion-exclusion

Require: Graph $G = (V, E)$ with $|V| = n$
Ensure: Chromatic number $\chi(G)$

```
F ← EnumerateIndepSets(G)
for k ← 1 to n do
    cnt ← 0
    for each S ⊆ V (represented as a bitset) do
        a ← |{I ∈ F : I & S = 0}|                    ▷ count IS avoiding S
        cnt ← cnt + (−1)^|S| · a^k
    end for
    if cnt ≠ 0 then return k
    end if
end for
```

4 IndSet-Removal Coloring

4.1 Independent Set Extraction

The IndSet extraction procedure (Algorithm 3) builds one independent set by iterating through the active vertices in index order. For each vertex v selected, it removes v and all of v's current neighbors from the active set.

Algorithm 3 Independent set extraction

Require: Graph $G = (V, E)$; active set $U \subseteq V$
Ensure: An independent set $I \subseteq U$

```
I ← ∅
while G has edges do
    v ← first vertex in U
    I ← I ∪ {v}
    Remove v and all neighbors of v from U and from G
end whilereturn I
```

4.2 IndSet-Removal Coloring

Algorithm 4 repeatedly calls the IS extractor to peel off one independent set (color class) at a time.

4.3 Quality and Complexity

IndSet-Removal Coloring produces at most $\Delta(G)+1$ colors, where $\Delta(G)$ is the maximum degree [2]. The vertex-order traversal makes it exact on complete graphs and odd cycles but can be arbitrarily far from optimal on bipartite or dense multipartite graphs (see Table 1 for empirical comparisons). Its primary advantage is $O(n+m)$ running time per color class, making it practical for large sparse graphs.

Algorithm 4 IndSet-Removal Coloring

Require: Graph $G = (V, E)$
Ensure: Coloring array $col[0 \ldots n-1]$

```
col ← array of size n; k ← 0
while G has edges do
    I ← ExtractIndepSet(G)
    for each v ∈ I do
        col[v] ← k
        Remove all edges incident to v from G
    end for
    k ← k + 1
end whilereturn col
```

5 Implementation

The project is written in C++17 and uses the Boost Graph Library (BGL) version 1.60 or later. The graph type is

```
typedef boost::adjacency_list<
    boost::vecS, boost::vecS, boost::undirectedS,
    property<boost::vertex_color_t, int>,
    property<boost::edge_weight_t, int,
             property<boost::edge_name_t, string>>> Graph;
```

Independent sets are stored as `boost::dynamic_bitset<>` objects, which support $O(n/w)$ bitwise AND, OR, and population-count operations where w is the machine word size.

The key functions are:

- `power_set(n, f)`: iterates over all 2^n subsets, calling `f` with each subset represented as a `dynbit`.
- `is_indep_set(G, B)`: checks all pairs of set bits in B for adjacency.
- `gen_ind_set_for_g(G)`: returns $\mathcal{F}(G)$ as a `std::set<dynbit>`.
- `alpha(S, F)`: counts $|\{I \in F : I \,\&\, S = 0\}|$ by a linear scan.
- `c_k(F, G, k)`: evaluates formula (1) using floating-point arithmetic and casts the result to `int`.

Generators for Cycle (C_n) and Generalized Peterson ($GP(n, k)$) graphs extend the abstract base class `GeneratorInterface` and override `generate()`.

Generalized Peterson graph. $GP(n, k)$ has $2n$ vertices. The outer n-cycle connects vertex i to $(i + 1) \bmod n$. Spoke edges connect outer vertex i to inner vertex $i + n$. Inner edges connect $i + n$ to $(i + n + k) \bmod n + n$ with a boundary correction.

Build system. The project uses CMake 3.10 or later. Three targets are built: `indset_exact_coloring` (IndSet-Based Exact Coloring demo, `demo_indset_exact_coloring.cpp`), `indset_removal_coloring` (IndSet-Removal demo, `demo_indset_removal_coloring.cpp`), and `tests` (`tests.cpp`, the test suite described in Section 6).

6 Experimental Results

All experiments were run on the `tests` binary compiled with `-std=c++17 --openmp -O2` on Ubuntu 22.04. Table 1 summarizes results for both the IndSet-Based Exact Coloring and IndSet-Removal Coloring algorithms.

All 78 test assertions pass: IndSet-Based Exact Coloring recovers the known chromatic number in every case. In particular, $c_k(K_9, 9) = 9! = 362{,}880$ and $c_k(K_{10}, 10) = 10! = 3{,}628{,}800$ are verified exactly. IndSet-Removal Coloring matches the exact result on complete graphs and odd cycles, but over-estimates on even cycles, bipartite, and multipartite graphs.

6.1 Exact Values of c_k at the Chromatic Number

For complete graphs, the formula (1) coincides with the standard chromatic polynomial at $k = \chi(G)$:

The coincidence for K_n can be verified algebraically: since the only independent sets of K_n are $\emptyset$ and the n singletons, $\alpha(S, \mathcal{F}) = n - |S| + 1$ and the sum (1) at $k = n$ evaluates to

$$\sum_{j=0}^{n} \binom{n}{j} (-1)^{n-j} (j+1)^n = n!,$$

which equals $P(K_n, n) = n!$. For cycle C_4 the coincidence is accidental; for C_5 the values differ.

6.2 Parity Property of Odd Cycles

For any odd cycle C_n (n odd) the algorithm correctly returns $\chi = 3$. No odd cycle is 2-colorable because any 2-coloring of a path forces the two endpoints of a cycle to receive the same color, contradicting the closing edge. For $n = 3, 5, 7$ the algorithm confirmed $c_k(C_n, 2) = 0$ and $c_k(C_n, 3) > 0$.

6.3 Petersen Graph

The Petersen graph $GP(5, 2)$ is the canonical example of a 3-chromatic 3-regular graph. It is not planar and has no Hamiltonian cycle. Our algorithm confirmed $\chi(GP(5, 2)) = 3$, with $c_k(GP(5, 2), 2) = 0$ and $c_k(GP(5, 2), 3) > 0$.

6.4 Wheel Graphs

The wheel graph W_n consists of a single hub vertex connected to every vertex of a cycle C_{n-1}. It has n vertices and $2(n-1)$ edges. Because the hub is adjacent to every rim vertex, the hub's color must differ from all colors used on the rim.

- If n is odd, the rim C_{n-1} has an even number of vertices and is therefore bipartite, so $\chi(C_{n-1}) = 2$. The hub requires a third color, giving $\chi(W_n) = 3$.
- If n is even, the rim C_{n-1} is an odd cycle, so $\chi(C_{n-1}) = 3$. The hub must avoid all three rim colors, giving $\chi(W_n) = 4$.

This makes wheel graphs the simplest natural family with $\chi = 4$. We verified $\chi(W_n)$ for $n \in \{5, 6, 7, 8, 9\}$; all 5 tests pass.

6.5 Complete Multipartite Graphs

The complete k-partite graph $K_{s,\ldots,s}$ (with k parts of s vertices each) has $\chi = k$: color each part with its own color (lower bound from the clique-like independence structure; k colors suffice since each part is an independent set). We tested:

- $K_{2,2,2}$ (octahedron): 6 vertices, 12 edges, $\chi = 3$.
- $K_{3,3,3}$: 9 vertices, 27 edges, $\chi = 3$.
- $K_{2,2,2,2}$: 8 vertices, 24 edges, $\chi = 4$.

These graphs have few independent sets relative to 2^n (only subsets within a single part), so the algorithm is fast despite the large edge count.

6.6 Generalized Petersen Family

Beyond $GP(5, 2)$ we tested five more members:

- $GP(4, 1) \cong Q_3$ (3-cube): bipartite, $\chi = 2$.
- $GP(5, 1)$ (pentagonal prism): odd outer cycle, $\chi = 3$.
- $GP(6, 1)$ (hexagonal prism): even outer cycle, bipartite, $\chi = 2$.
- $GP(6, 2)$: inner triangles force $\chi \geq 3$; algorithm confirms $\chi = 3$.
- $GP(7, 2)$: odd outer and inner cycles, $\chi = 3$.

6.7 SuiteSparse Matrix Collection Graphs

To demonstrate applicability to real-world data, we obtained matrices from the SuiteSparse Matrix Collection [6] and converted each to an undirected graph by symmetrizing (adding edge (i, j) whenever $A_{ij} \neq 0$ or $A_{ji} \neq 0$) and discarding self-loops and edge weights. For matrices larger than 14 vertices we extract a connected subgraph of k vertices via BFS from a fixed starting vertex.

The cage matrices arise from random walks on finite words over the DNA alphabet $\{A, C, G, T\}$; `cage3` has $|\Sigma|^3 - |\Sigma|^2 = 48$ directed edges reduced to 7 undirected after symmetrization. Both DNA graphs require exactly 3 colors — they contain odd cycles formed by overlapping DNA words that share a common suffix/prefix. The `GD98_a` subgraph also requires 3 colors, while the symmetric FEM matrix `ndtest` is 2-colorable in its 10-vertex neighborhood, consistent with the bipartite structure typical of finite-element adjacency graphs.

6.8 Scaling Analysis

Table 4 reports wall-clock times for selected larger graphs. The dominant cost is evaluating c_k (formula (1)), which is $O(2^n \cdot |\mathcal{F}|)$ per candidate k. For sparse graphs such as cycles, $|\mathcal{F}|$ grows like the Lucas sequence $L(n) \approx \varphi^n$, so the total cost is roughly $O((2\varphi)^n) \approx O(3.24^n)$, yielding approximately a 10× slowdown per 2 additional vertices. Dense graphs such as $K_{m,m}$ have much smaller $|\mathcal{F}|$ and therefore scale better.

6.9 High Chromatic Number Graphs

Complete graphs K_n have $\chi(K_n) = n$ and $|\mathcal{F}| = n + 1$ (only singletons and the empty set), so the independent-set enumeration phase is trivial. The bottleneck is computing the chromatic polynomial at each candidate k, which requires intermediate terms of the form $(-1)^{|S|}\,\alpha(S, \mathcal{F})^k$. For K_9, these terms reach $9^9 \approx 3.9 \times 10^8$, which fits comfortably in a 64-bit integer. The implementation uses `long long` accumulation with integer exponentiation (no floating point), giving correct results up to at least K_{10} ($10! = 3{,}628{,}800$, verified by the test suite).

Complete k-partite graphs $K_{2\times k}$ (two vertices per part) provide compact benchmarks for high chromatic numbers: only $2k$ vertices but $\chi = k$. Their independent-set families have exactly 3^k members (each part contributes three choices: both, neither, or one vertex), far fewer than 2^{2k}, so the algorithm scales well despite the large chromatic number.

6.10 IndSet-Removal vs. Exact Coloring

The IndSet-Removal column in Tables 1–5 shows the number of colors produced by Algorithm 4 on each graph. The algorithm processes vertices in ascending index order, so its quality depends strongly on the vertex ordering. The results illustrate several patterns:

- **Complete graphs**: IndSet-Removal is always optimal. Because every vertex is adjacent to all lower-indexed vertices, each vertex receives a new color, achieving exactly n colors for K_n.
- **Odd cycles**: IndSet-Removal is optimal ($\chi = 3$). Vertex 0 and its non-adjacent neighbor both receive color 0, and the single remaining vertex requires a third color.
- **Even cycles and bipartite graphs**: IndSet-Removal performs poorly because the vertex-order traversal fails to exploit the bipartite structure. C_4 receives 4 colors instead of 2; similarly $GP(6, 1)$ (hexagonal prism, bipartite) receives 6 instead of 2.
- **Wheel graphs**: exact when n is even (odd rim forces 4 colors regardless of ordering) but over-estimates by 2 when n is odd.
- **Complete multipartite** $K_{2\times k}$: the over-count grows linearly with k; IndSet-Removal uses approximately $2k - 1$ colors instead of k, giving an approximation ratio approaching 2 as k grows.

These results confirm that IndSet-Removal Coloring is a fast but unreliable upper bound: it is tight on graphs where the natural vertex order aligns with the chromatic structure (complete graphs, odd cycles), but can be far from optimal on structured sparse or multipartite graphs.

7 Complexity Analysis

7.1 Independent Set Enumeration

`gen_ind_set_for_g` calls `power_set`, which iterates over all 2^n subsets. For each subset B, `is_indep_set` checks all pairs of set bits: at most $O(n^2)$ edge queries. Hence the enumeration phase runs in $O(2^n \cdot n^2)$ time and produces a set $\mathcal{F}$ of at most $3^n/\mathrm{poly}(n)$ independent sets (by the Moon–Moser theorem on the maximum number of maximal independent sets in a graph with n vertices [3]).

7.2 Evaluation of c_k

For each candidate k, evaluating formula (1) requires iterating over all 2^n subsets S and, for each S, scanning all $|\mathcal{F}|$ independent sets to compute $\alpha(S, \mathcal{F})$. This costs $O(2^n \cdot |\mathcal{F}| \cdot n/w)$ per value of k, where n/w is the cost of a bitwise AND. Summing over at most n values of k, the total evaluation cost is $O(n \cdot 2^n \cdot |\mathcal{F}| \cdot n/w)$.

7.3 Overall Complexity

The dominating term is the evaluation of c_k, giving

$$T(n) = O\big(n^2 \cdot 2^n \cdot |\mathcal{F}(G)|\big)\,.$$

In the worst case (e.g., the empty graph E_n has $|\mathcal{F}| = 2^n$), this is $O(n^2 \cdot 4^n)$. Empirically (Table 4), sparse graphs such as cycles reach the practical limit around $n = 13$–14 (seconds to minutes), while dense graphs such as $K_{m,m}$ scale better because $|\mathcal{F}|$ is much smaller than 2^n. The algorithm is intended for IndSet-Based Exact Coloring of small subgraphs; for large graphs the IndSet-Removal coloring in `demo_is_removal_coloring.cpp` is the appropriate choice.

7.4 Space Complexity

Storing $\mathcal{F}$ requires $O(|\mathcal{F}| \cdot n/w)$ words, which is $O(3^n \cdot n/w)$ in the worst case.

8 Conclusion

We have described and formally justified the inclusion-exclusion algorithm for computing the chromatic number of a graph implemented in the `graph-exact-coloring-heuristic` project. The key result is that the quantity $c_k(G, k)$ — the number of ordered k-tuples of independent sets covering V — is zero if and only if G is not k-colorable (Theorem 2). Scanning $k = 1, 2, \ldots$ and stopping at the first nonzero value thus yields $\chi(G)$ exactly.

The algorithm is exponential in $|V|$ and is therefore intended for small graphs or for IndSet-Based Exact Coloring of subgraphs within a larger pipeline. The companion IndSet-Removal Coloring algorithm runs in $O(n + m)$ time per color class and provides a fast upper bound useful as an initialization strategy or a standalone approximation for large instances.

Future directions include: (i) exploiting symmetry to prune the power-set enumeration, (ii) integrating branch-and-bound to improve worst-case behavior, and (iii) leveraging OpenMP (already supported in the build) to parallelize the inner α-counting loop.

References

[1] R. M. Karp. Reducibility among combinatorial problems. In R. E. Miller and J. W. Thatcher, editors, *Complexity of Computer Computations*, pages 85–103. Plenum Press, 1972.

[2] D. J. A. Welsh and M. B. Powell. An upper bound for the chromatic number of a graph and its application to timetabling problems. *The Computer Journal*, 10(1):85–86, 1967.

[3] J. W. Moon and L. Moser. On cliques in graphs. *Israel Journal of Mathematics*, 3(1):23–28, 1965.

[4] E. L. Lawler. A note on the complexity of the chromatic number problem. *Information Processing Letters*, 5(3):66–67, 1976.

[5] J. G. Siek, L.-Q. Lee, and A. Lumsdaine. *The Boost Graph Library: User Guide and Reference Manual*. Addison-Wesley, 2002.

[6] T. A. Davis and Y. Hu. The University of Florida sparse matrix collection. *ACM Transactions on Mathematical Software*, 38(1):1–25, 2011.

Table 1: All three results per graph: known χ, IndSet-Based Exact Coloring, and IndSet-Removal Coloring (78 exact assertions, all passing).

Graph	Description	$\lvert V\rvert$	$\lvert E\rvert$	Known χ	IndSet-Exact	IndSet-Removal
Cycle graphs						
C_3	Triangle	3	3	3	3	3
C_4	4-cycle (square)	4	4	2	2	4
C_5	5-cycle (pentagon)	5	5	3	3	3
C_6	6-cycle (hexagon)	6	6	2	2	4
C_7	7-cycle	7	7	3	3	3
Complete and edgeless graphs						
K_1	Single vertex	1	0	1	1	1
K_2	Single edge	2	1	2	2	2
K_3	Triangle	3	3	3	3	3
K_4	Complete graph	4	6	4	4	4
K_9	Complete graph	9	36	9	9	9
K_{10}	Complete graph	10	45	10	10	10
E_1	Isolated vertex	1	0	1	1	1
E_3	Three isolated vertices	3	0	1	1	1
Bipartite graphs						
$K_{2,3}$	Complete bipartite	5	6	2	2	3
Wheel graphs W_n; $\chi = 3$ if n odd, $\chi = 4$ if n even						
W_5	Hub $+ C_4$ (even rim)	5	8	3	3	5
W_6	Hub $+ C_5$ (odd rim)	6	10	4	4	4
W_7	Hub $+ C_6$ (even rim)	7	12	3	3	5
W_8	Hub $+ C_7$ (odd rim)	8	14	4	4	4
W_9	Hub $+ C_8$ (even rim)	9	16	3	3	5
Complete multipartite graphs; $\chi(K^{(k)}_{s,\ldots,s}) = k$						
$K_{2,2,2}$	Octahedron	6	12	3	3	5
$K_{3,3,3}$	Complete 3-partite	9	27	3	3	7
$K_{2,2,2,2}$	Complete 4-partite	8	24	4	4	7
$K_{2\times 5}$	Complete 5-partite	10	40	5	5	9
$K_{2\times 6}$	Complete 6-partite	12	60	6	6	11
$K_{2\times 7}$	Complete 7-partite	14	84	7	7	13
$K_{2\times 8}$	Complete 8-partite	16	112	8	8	15
$K_{2\times 9}$	Complete 9-partite	18	144	9	9	17
Generalized Petersen graphs $GP(n,k)$						
$GP(4,1)$	3-cube Q_3 (bipartite)	8	12	2	2	6
$GP(5,1)$	Pentagonal prism	10	15	3	3	3
$GP(5,2)$	Petersen graph	10	15	3	3	4
$GP(6,1)$	Hexagonal prism	12	18	2	2	6
$GP(6,2)$	—	12	18	3	3	4
$GP(7,2)$	—	14	21	3	3	4

Table 2: Values of $c_k(G, \chi(G))$ vs. the chromatic polynomial $P(G, \chi(G))$.

Graph	χ	$c_k(G, \chi)$	$P(G, \chi)$	Match?
K_1	1	1	1	Yes
K_2	2	2	2	Yes
K_3	3	6	6	Yes
K_4	4	24	24	Yes
C_3	3	6	6	Yes
C_4	2	2	2	Yes
C_5	3	60	30	No

Table 3: SuiteSparse graphs: known χ, IndSet-Based Exact Coloring, IndSet-Removal Coloring, and exact algorithm wall-clock time.

Matrix	Group	Processing	$\lvert V\rvert$	$\lvert E\rvert$	Known χ	IndSet-Exact	IndSet-Removal	Time
`cage3`	vanHeukelum	symmetrized (full)	5	7	3	3	4	<1 ms
`cage4`	vanHeukelum	symmetrized (full)	9	20	3	3	7	3 ms
`GD98_a`	Pajek	BFS-10 subgraph	10	10	3	3	3	81 ms
`ndtest`	(project)	BFS-10 subgraph	10	14	2	2	6	15 ms

Table 4: Wall-clock time for larger graphs (Release build, single core, Ubuntu 22.04). IndSet-Exact gives the exact chromatic number; IndSet-Removal is the heuristic upper bound.

Graph	Family	$\lvert V\rvert$	$\lvert E\rvert$	Known χ	IndSet-Exact	IndSet-Removal	Time
C_8	Cycle	8	8	2	2	4	1 ms
C_9	Cycle	9	9	3	3	3	7 ms
C_{10}	Cycle	10	10	2	2	4	24 ms
C_{11}	Cycle	11	11	3	3	3	86 ms
C_{12}	Cycle	12	12	2	2	4	168 ms
C_{13}	Cycle	13	13	3	3	3	786 ms
C_{14}	Cycle	14	14	2	2	4	≈3 s
C_{16}	Cycle	16	16	2	2	4	≈33 s
K_5	Complete	5	10	5	5	5	<1 ms
K_6	Complete	6	15	6	6	6	<1 ms
K_7	Complete	7	21	7	7	7	<1 ms
K_8	Complete	8	28	8	8	8	1 ms
$K_{4,4}$	Bipartite	8	16	2	2	5	1 ms
$K_{5,5}$	Bipartite	10	25	2	2	6	9 ms
$K_{6,6}$	Bipartite	12	36	2	2	7	69 ms
$K_{7,7}$	Bipartite	14	49	2	2	8	536 ms
$GP(6,2)$	Gen. Petersen	12	18	3	3	4	135 ms
$GP(7,2)$	Gen. Petersen	14	21	3	3	4	1333 ms
$GP(8,3)$	Gen. Petersen	16	24	2	2	5	≈15 s

Table 5: High chromatic number graphs: known χ, IndSet-Based Exact Coloring, IndSet-Removal Coloring, and exact algorithm wall-clock time.

Graph	V	E	Known χ	IndSet-Exact	IndSet-Removal	Time
K_9	9	36	9	9	9	3 ms
K_{10}	10	45	10	10	10	9 ms
$K_{2\times5}$	10	40	5	5	9	7 ms
$K_{2\times6}$	12	60	6	6	11	47 ms
$K_{2\times7}$	14	84	7	7	13	287 ms
$K_{2\times8}$	16	112	8	8	15	1.1 s
$K_{2\times9}$	18	144	9	9	17	5.5 s